
Rapport AIT - Lab04 - Docker

Nathan Séville, Julien Quartier

Table des matières

Introduction	3
Tasks	3
Task 0: Identify issues and install the tools	3
Task 1: Add a process supervisor to run several processes	4
Task 2: Add a tool to manage membership in the web server cluster	5
Difficulties	5
Conclusion	5

Introduction

Tasks

Task 0: Identify issues and install the tools

[M1] Do you think we can use the current solution for a production environment? What are the main problems when deploying it in a production environment?

Non, cette solution n'est pas adaptée à un environnement de production. En cas d'arrêt inopiné de *node*, aucun monitoring, ni procédure automatique n'est configurée. En cas de grande charge, aucune stratégie de *scaling* n'est définie. L'ajout de nouveau *container* est compliquée dans l'infrastructure courante (CF. M2).

[M2] Describe what you need to do to add new *webapp* container to the infrastructure. Give the exact steps of what you have to do without modifying the way the things are done. Hint: You probably have to modify some configuration and script files in a Docker image.

///// To complete

1. Ajouter une *webapp* dans le fichier `docker-compose.yml`.
2. Ajouter une *node* dans le fichier de configuration de *haproxy*, `haproxy.cfg`.

[M3] Based on your previous answers, you have detected some issues in the current solution. Now propose a better approach at a high level.

[M4] You probably noticed that the list of web application nodes is hardcoded in the load balancer configuration. How can we manage the web app nodes in a more dynamic fashion?

[M5] In the physical or virtual machines of a typical infrastructure we tend to have not only one main process (like the web server or the load balancer) running, but a few additional processes on the side to perform management tasks.

For example to monitor the distributed system as a whole it is common to collect in one centralized place all the logs produced by the different machines. Therefore we need a process running on each machine that will forward the logs to the central place. (We could also imagine a central tool that reaches out to each machine to gather the logs. That's a push vs. pull problem.) It is quite common to see a push mechanism used for this kind of task.

Do you think our current solution is able to run additional management processes beside the

2. Un process supervisor est important dans une infrastructure load balancée pour pouvoir manager nos *nodes* en fonction de la charge ainsi que de relancer les *nodes* qui peuvent possiblement “crasher”. Il peut s’avérer utile de supprimer les *nodes* non utilisées d’autant plus si la stratégie du *load balancer* et *first* par exemple.

// TO COMPLETE

Task 2: Add a tool to manage membership in the web server cluster

Difficulties

Conclusion